



# Debezium

---

**CDC - CHANGE DATA CAPTURE**

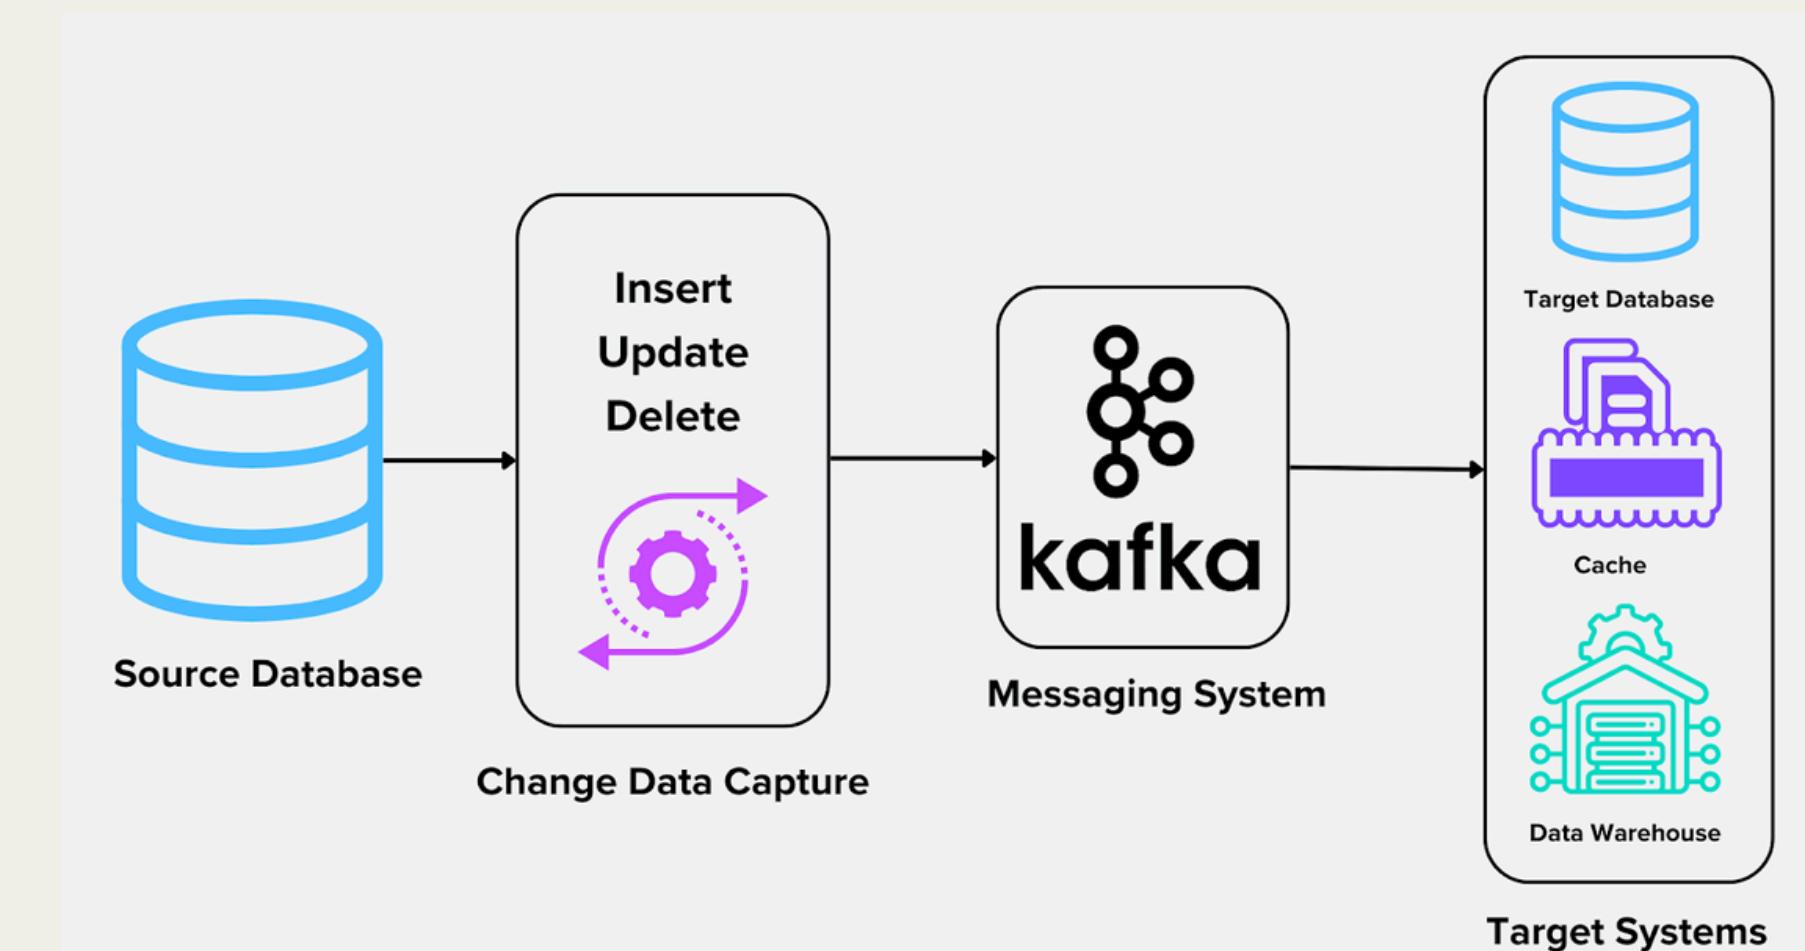
# SADRŽAJ

---



- Šta je CDC
- Tipični CDC use cases
- Šta je Debezium
- Konkurentni pristupi CDC-u
- Kako Debezium implementira CDC
- Debezium Source Connectors
  - Snapshotting
  - Streaming
- Debezium Sink Connectors
- Custom Sink Connectors
- Projekti

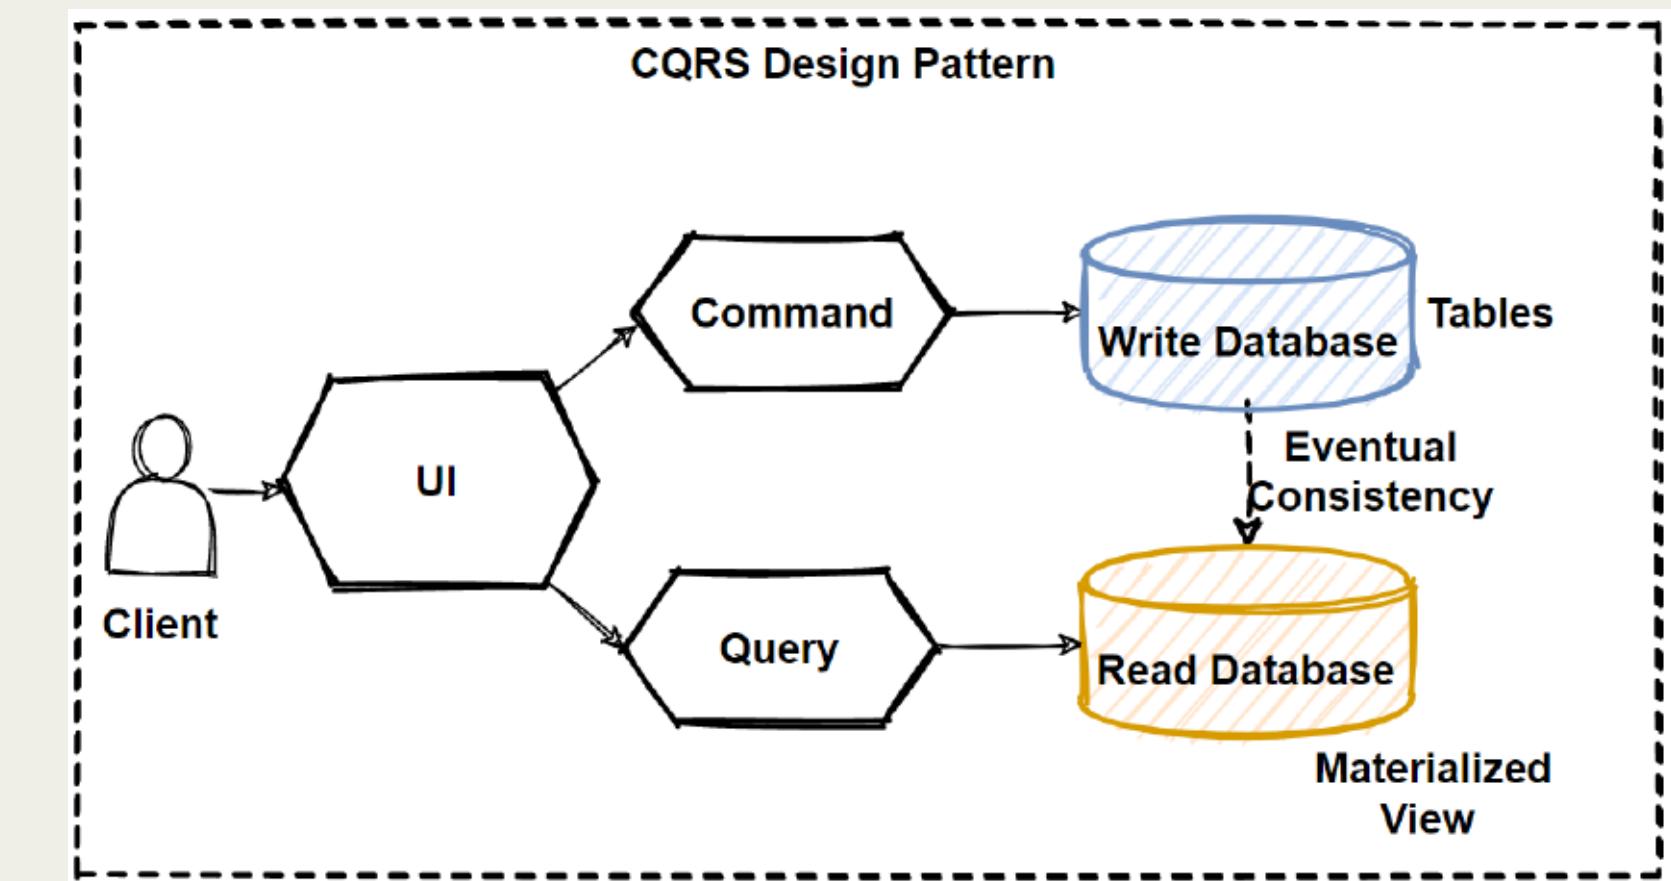
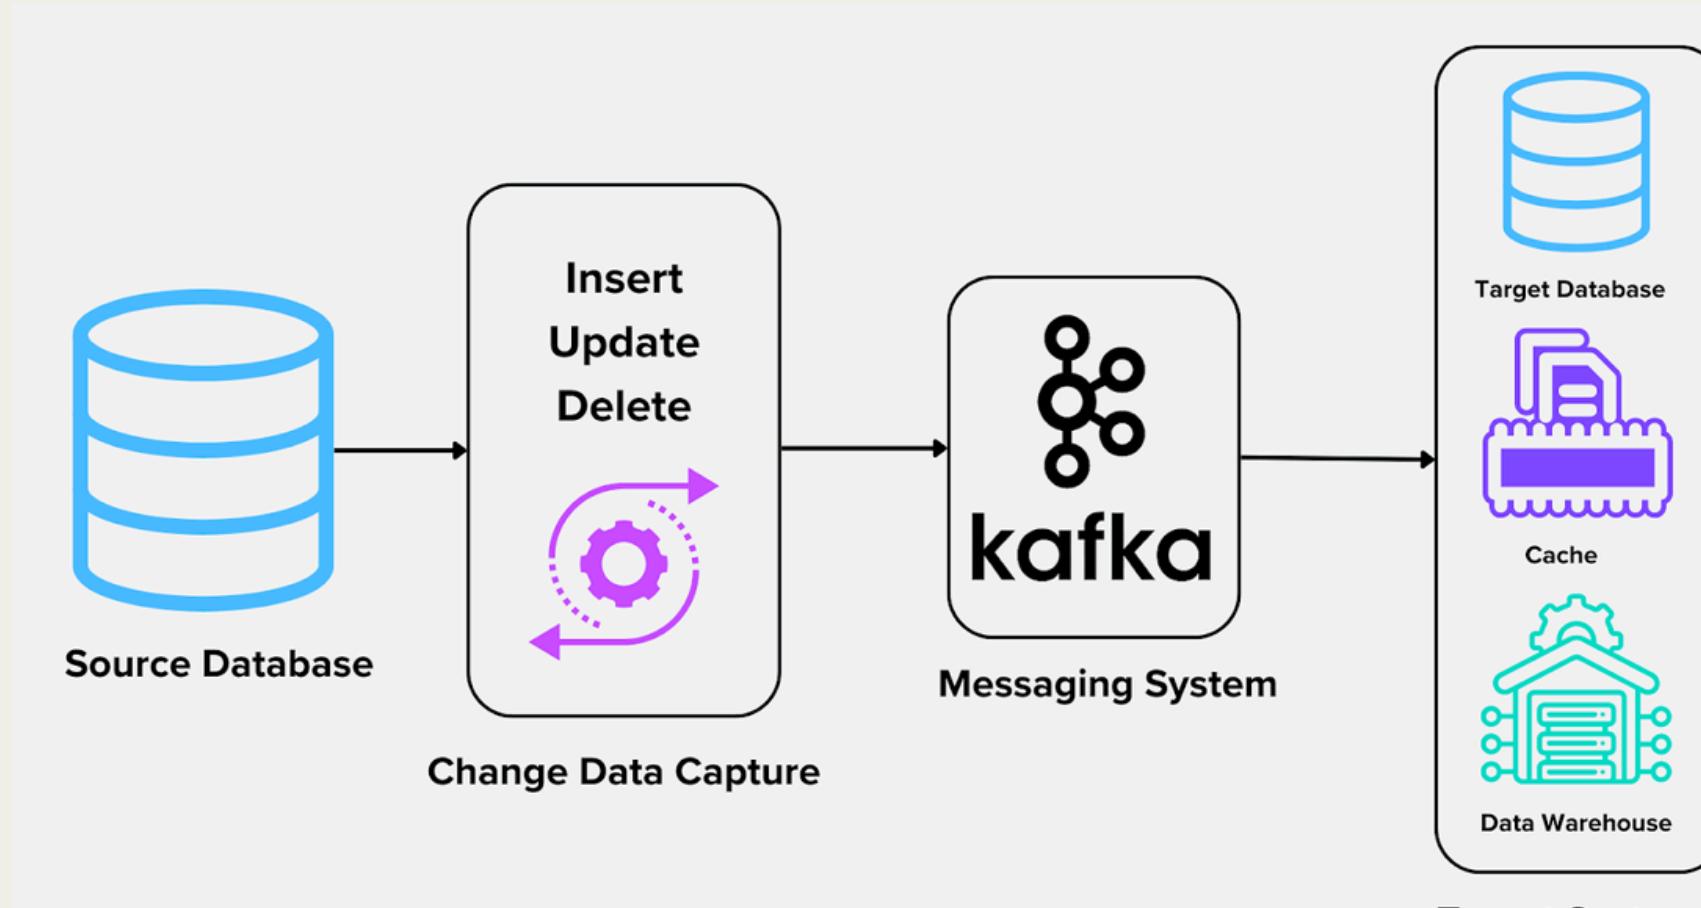
- Tehnika koja omogućava praćenje promena nad podacima u bazi podataka
- Najčešće podrazumeva da se promene nad podacima objavljaju korišćenjem nekog sistema za razmenu poruka (najčešće Kafka)



# TIPičNI CDC USE CASES



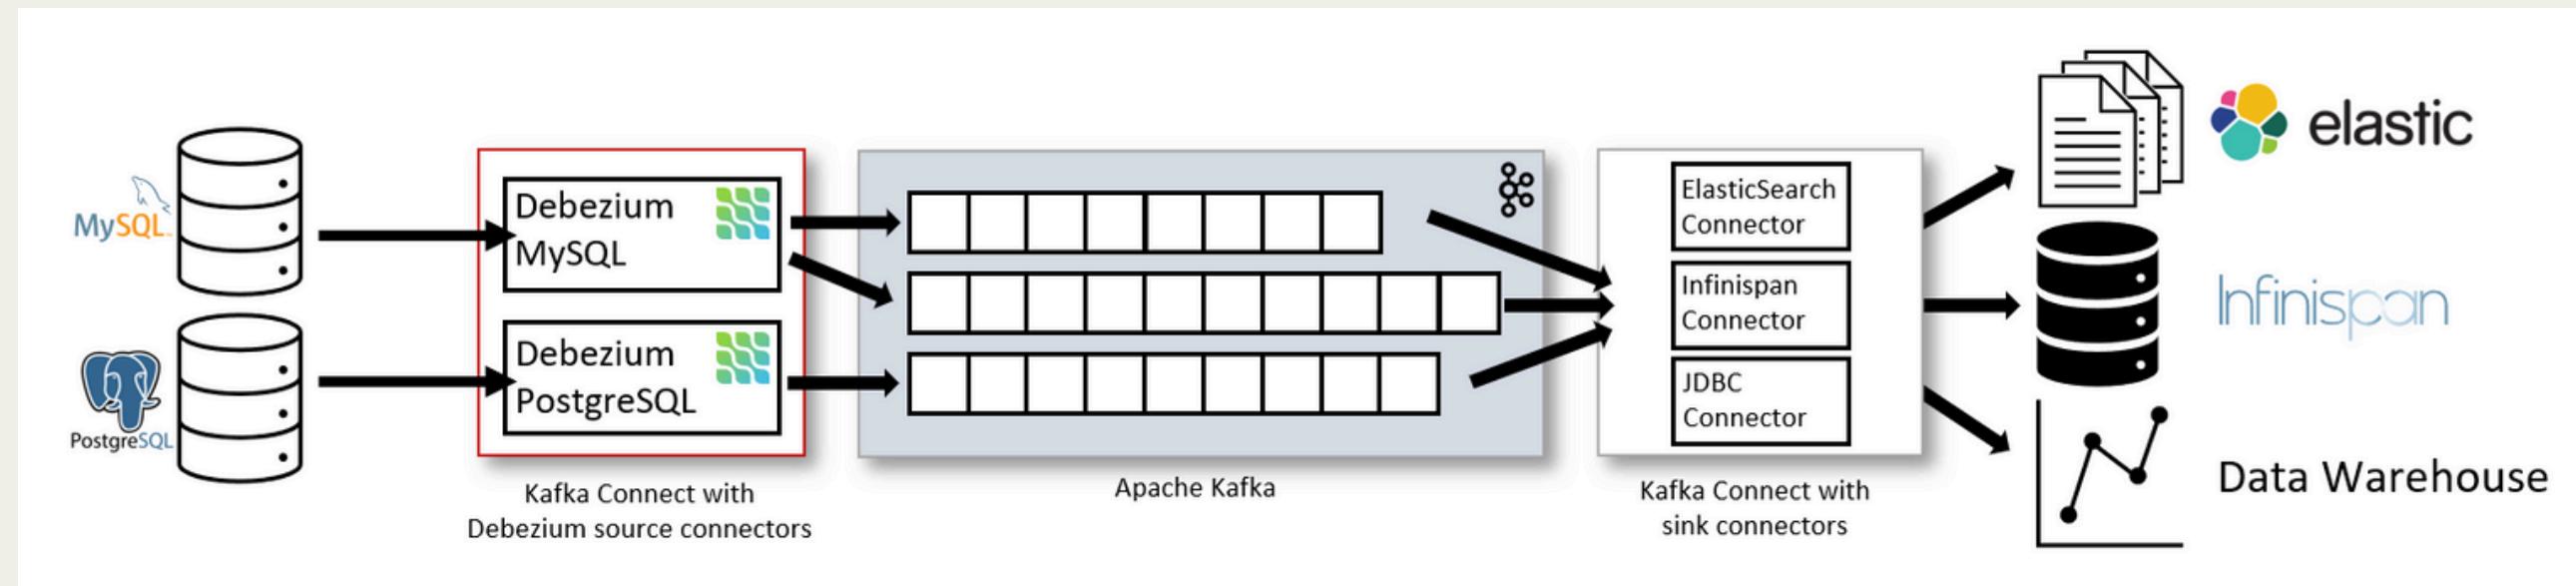
- Replikacija podataka
  - Upis u veći broj baza podataka
  - Održavanje keš koherencije
  - CQRS pattern
- } EVENTUAL CONSISTENCY



# ŠTA JE DEBEZIUM



- Open Source platforma koja se koristi za implementaciju CDC-a
- Koristi Apache Kafku
- Koristi KafkaConnect
- Source Connectors - povezuje se na izvor i streamuje promene nad podacima
  - Podrška za PostgreSQL, Oracle, MySQL, MariaDB, SQL Server, Db2, Cassandra...
- Sink Connectors - konzumira evente sa Kafke i upisuje u odredište
  - Podrška za bilo koju bazu za koju postoji JDBC plugin, MongoDB



- Trigger based CDC

```
CREATE TABLE users_cdc (
    cdc_id BIGSERIAL PRIMARY KEY,
    operation CHAR(1), -- C, U, D
    customer_id INT,
    old_data JSONB,
    new_data JSONB,
    changed_at TIMESTAMP DEFAULT now()
);
```

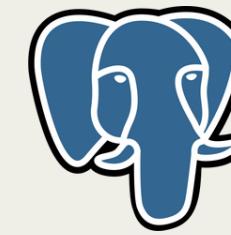
```
CREATE TRIGGER users_cdc_trg
AFTER INSERT OR UPDATE OR DELETE
ON users
FOR EACH ROW
EXECUTE FUNCTION users_cdc_trigger();
```

```
CREATE OR REPLACE FUNCTION users_cdc_trigger()
RETURNS TRIGGER AS $$
BEGIN
    IF TG_OP = 'INSERT' THEN
        INSERT INTO users_cdc (operation, customer_id, old_data, new_data)
        VALUES ('C', NEW.id, NULL, to_jsonb(NEW));
        RETURN NEW;
    ELSIF TG_OP = 'UPDATE' THEN
        INSERT INTO users_cdc (operation, customer_id, old_data, new_data)
        VALUES ('U', NEW.id, to_jsonb(OLD), to_jsonb(NEW));
        RETURN NEW;
    ELSIF TG_OP = 'DELETE' THEN
        INSERT INTO users_cdc (operation, customer_id, old_data, new_data)
        VALUES ('D', OLD.id, to_jsonb(OLD), NULL);
        RETURN OLD;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

- Log based CDC
- Specijalni fajlovi u koje baza beleži sve promene koje se dešavaju nad podacima
- Baza prvo upisuje promenu u log, a tek onda u samu tabelu
- Osnovna namena - oporavak, A i D iz ACID
  - PostgreSQL - Write Ahead Log (WAL)
  - MySQL - Binary Log (binlog)
  - Oracle - Redo Log
  - SQL Server - Transaction Log

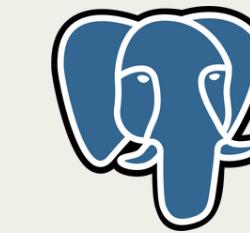
# DEBEZIUM SOURCE CONNECTORS - SNAPSHOTTING

---



- Radi se nakon uspostavljanja inicijalne konekcije (konfigurabilno)
  - Pokreće se read only transakcija
  - Čitaju se svi podaci iz odgovarajućih tabela
  - Za svaki pročitani red šalje se event na Kafku
- Ad hoc snapshots
  - Paralelni mod - streaming i snapshotting
- Inkrementalni snapshot
  - Korisnik ima mogućnost da pokrene snapshot u bilo kom trenutku

# DEBEZIUM SOURCE CONNECTORS - STREAMING



- Ime topic-a se po default-u formira kao <ime baze>.<ime šeme>.<tabela> (konfigurabilno)

- Sve promene na jedan topic:

```
"transforms": "Reroute",
"transforms.Reroute.topic.regex": "dbz\\\\.public\\..*",
"transforms.Reroute.topic.replacement": "CDC_1",
"transforms.Reroute.type": "io.debezium.transforms.ByLogicalTableRouter"
```

- Eventi mogu da se emituju u različitim formatima (JSON, Avro)

- before - vrednost koju je red imao pre (null, ako je operacija insert)
  - after - nova vrednost reda (null, ako je operacija delete)
  - source - izvor odakle dolazi event
  - op - operacija
    - r - read (u snapshot modu)
    - c - create
    - u - update
    - d - delete
  - message (postgres periodično u WAL upisuje stvari koje se ne odnose na podatke u bazi, message op se odnosi na ovaku aktivnost nad WAL-om)
  - ostali meta podaci, kao što je id transakcije koja je izvršila promenu, vreme izvršenja i slično.

# DEBEZIUM SOURCE CONNECTORS - STREAMING



```
{  
  "before": null,  
  "after": {  
    "user_id": 1,  
    "attribute_name": "height",  
    "attribute_value": "182"  
  },  
  "source": {  
    "version": "2.7.3.Final",  
    "connector": "postgresql",  
    "name": "dbz",  
    "ts_ms": 1767813361449,  
    "snapshot": "false",  
    "db": "postgres",  
    "sequence": "[\"25185064\", \"25185120\"]",  
    "ts_us": 1767813361449509,  
    "ts_ns": 1767813361449509000,  
    "schema": "public",  
    "table": "user_attributes",  
    "txId": 770,  
    "lsn": 25185120,  
    "xmin": null  
  },  
  "transaction": null,  
  "op": "c",  
  "ts_ms": 1767813361911,  
  "ts_us": 1767813361911197,  
  "ts_ns": 1767813361911197727  
}
```

```
{  
  "before": {  
    "user_id": 1,  
    "attribute_name": "height",  
    "attribute_value": "182"  
  },  
  "after": {  
    "user_id": 1,  
    "attribute_name": "height",  
    "attribute_value": "183"  
  },  
  "source": {  
    "version": "2.7.3.Final",  
    "connector": "postgresql",  
    "name": "dbz",  
    "ts_ms": 1767813524088,  
    "snapshot": "false",  
    "db": "postgres",  
    "sequence": "[\"25185520\", \"25186272\"]",  
    "ts_us": 1767813524088716,  
    "ts_ns": 1767813524088716000,  
    "schema": "public",  
    "table": "user_attributes",  
    "txId": 771,  
    "lsn": 25186272,  
    "xmin": null  
  },  
  "transaction": null,  
  "op": "u",  
  "ts_ms": 1767813524338,  
  "ts_us": 1767813524338060,  
  "ts_ns": 1767813524338060199  
}
```

```
{  
  "before": {  
    "user_id": 1,  
    "attribute_name": "height",  
    "attribute_value": "183"  
  },  
  "after": null,  
  "source": {  
    "version": "2.7.3.Final",  
    "connector": "postgresql",  
    "name": "dbz",  
    "ts_ms": 1767813555180,  
    "snapshot": "false",  
    "db": "postgres",  
    "sequence": "[\"25186424\", \"25186480\"]",  
    "ts_us": 1767813555180271,  
    "ts_ns": 1767813555180271000,  
    "schema": "public",  
    "table": "user_attributes",  
    "txId": 772,  
    "lsn": 25186480,  
    "xmin": null  
  },  
  "transaction": null,  
  "op": "d",  
  "ts_ms": 1767813555297,  
  "ts_us": 1767813555297637,  
  "ts_ns": 1767813555297637410  
}
```

# DEBEZIUM SOURCE CONNECTORS - STREAMING



```
{  
  "before": null,  
  "after": {  
    "user_id": 1,  
    "attribute_name": "height",  
    "attribute_value": "182"  
  },  
  "source": {  
    "version": "2.7.3.Final",  
    "connector": "postgresql",  
    "name": "dbz",  
    "ts_ms": 1767813361449,  
    "snapshot": "false",  
    "db": "postgres",  
    "sequence": "[\"25185064\", \"25185120\"]",  
    "ts_us": 1767813361449509,  
    "ts_ns": 1767813361449509000,  
    "schema": "public",  
    "table": "user_attributes",  
    "txId": 770,  
    "lsn": 25185120,  
    "xmin": null  
  },  
  "transaction": null,  
  "op": "c",  
  "ts_ms": 1767813361911,  
  "ts_us": 1767813361911197,  
  "ts_ns": 1767813361911197727  
}
```

```
{  
  "schema": {  
    "type": "struct",  
    "fields": [  
      {  
        "type": "struct",  
        "fields": [  
          {  
            "type": "int32",  
            "optional": false,  
            "field": "user_id"  
          },  
          {  
            "type": "string",  
            "optional": false,  
            "field": "attribute_name"  
          },  
          {  
            "type": "string",  
            "optional": true,  
            "field": "attribute_value"  
          }  
        ],  
        "optional": true,  
        "name": "CDC_1.Value",  
        "field": "before"  
      },  
      {  
        "type": "struct",  
        "fields": [  
          {  
            "type": "int32",  
            "optional": false,  
            "field": "user_id"  
          },  
          {  
            "type": "string",  
            "optional": false,  
            "field": "attribute_name"  
          },  
          {  
            "type": "string",  
            "optional": true,  
            "field": "attribute_value"  
          }  
        ],  
        "optional": true,  
        "name": "CDC_1.Value",  
        "field": "after"  
      },  
      {  
        "type": "struct",  
        "fields": [  
          {  
            "type": "string",  
            "optional": false,  
            "field": "version"  
          },  
          {  
            "type": "string",  
            "optional": false,  
            "field": "connector"  
          },  
          {  
            "type": "string",  
            "optional": false,  
            "field": "name"  
          }  
        ]  
      }  
    ]  
  }  
}
```

- KafkaConnect sink konektori, koji konzumiraju evente sa Kafka topic-a i upisuju u neki drugi sistem
  - Podrška za bilo koju bazu podataka za koju postoji JDBC driver
  - MongoDB
- JDBC konektori nude:
  - At-least-once delivery - svaki event koji se pričita sa Kafke, biće i obrađen
  - Automatsko mapiranje tipova podataka i tipova kolona
  - Idempotentni upisi - nije podrazumevano podešavanje, moguće je podesiti `insert.mode` atribut konektora
  - Schema evolution - osnovna podrška za detektovanje promena u šemi izvorene tabele

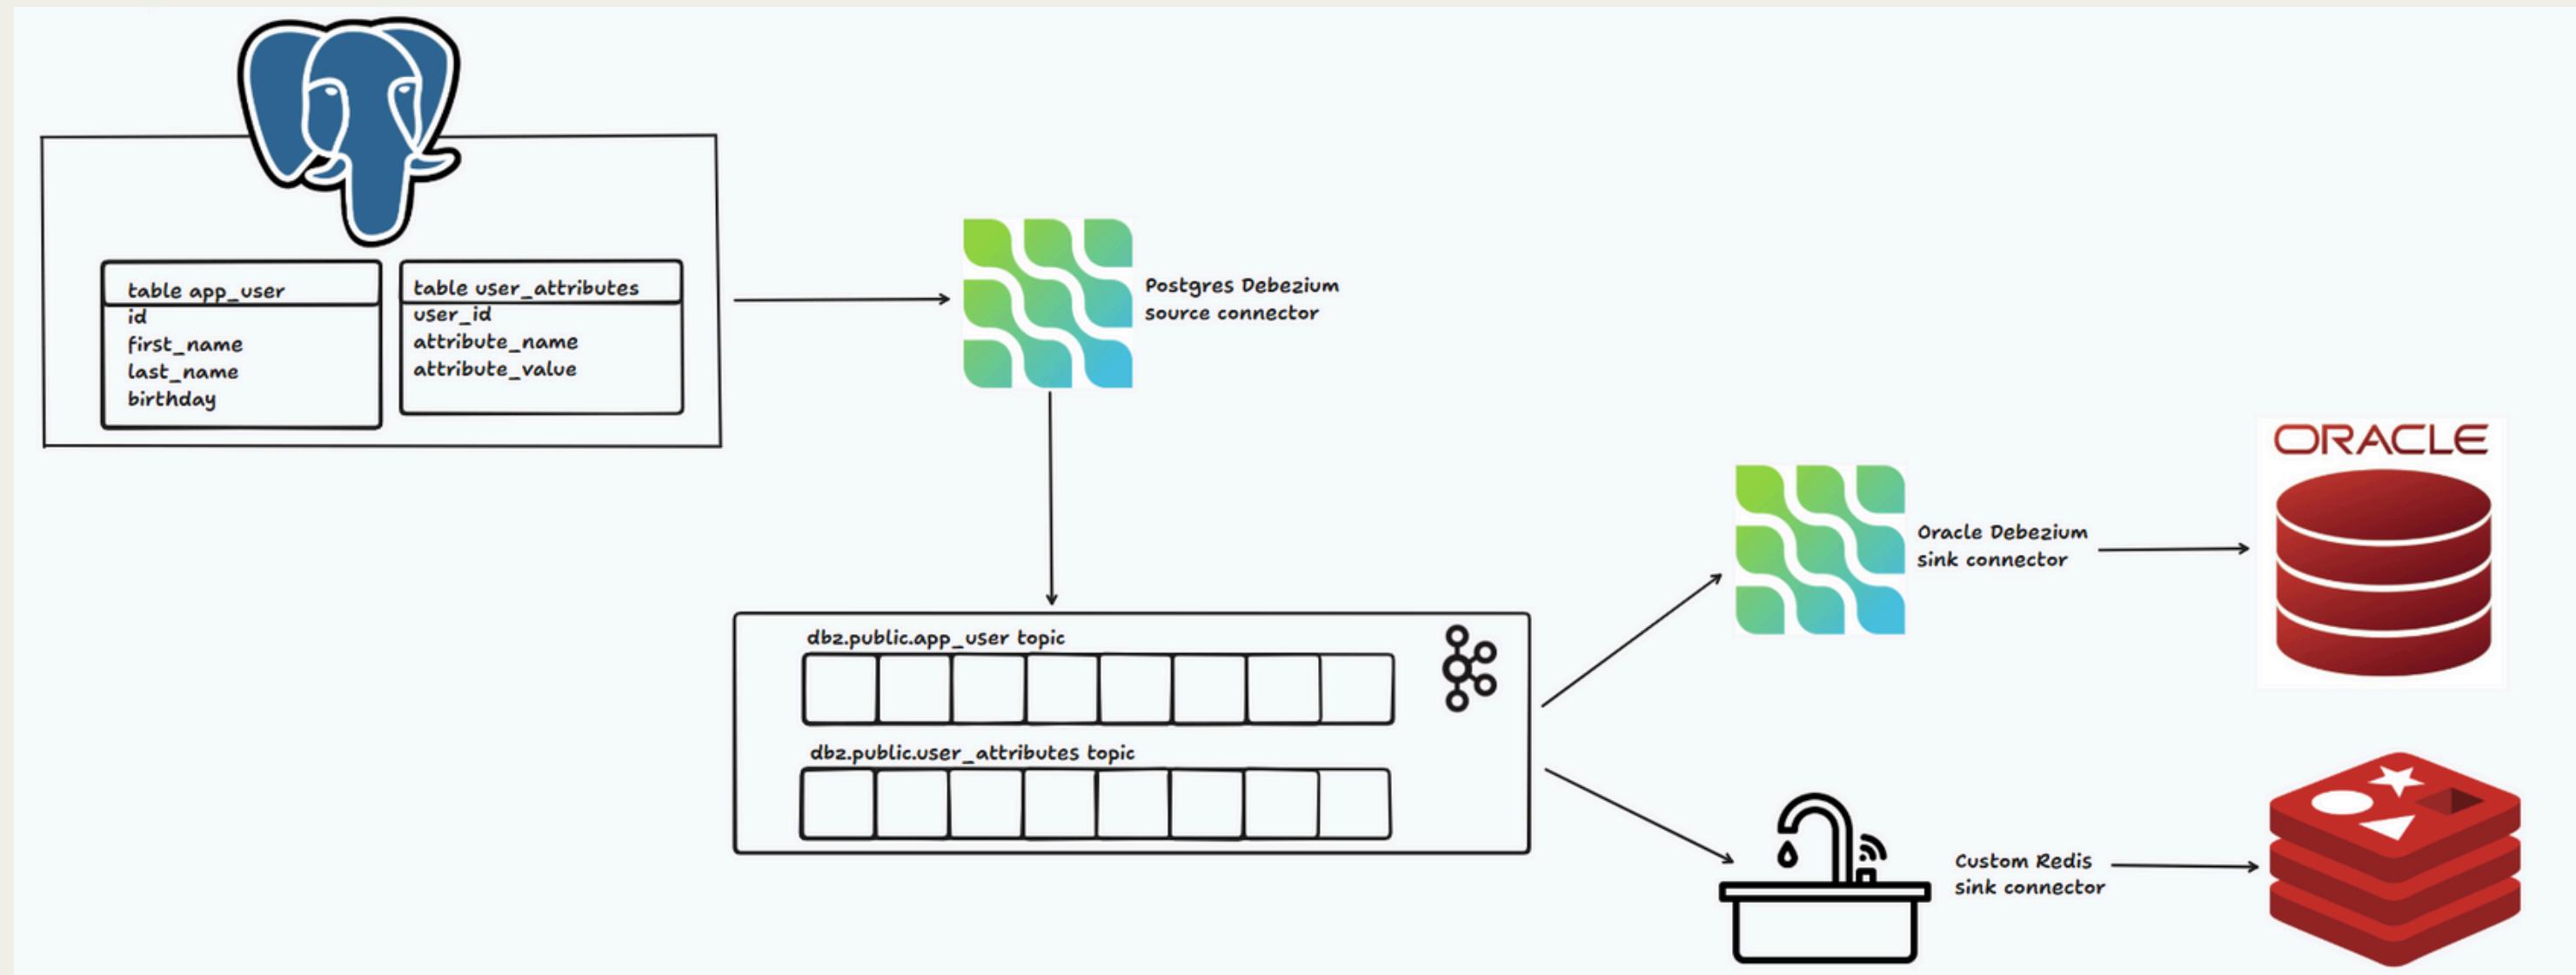
# CUSTOM SINK CONNECTORS

---



- Generalno, Debezium se najčešće koristi kako bi se promene nad bazom podataka emitovale na Kafku u obliku evenata
- Custom Sink Connectori nisu ništa više do obični Kafka consumer-i koji se povezuju na Kafku i mogu da vrše čitanje evenat-a
- Specijalizovani su, tako da se može izostaviti deo koji se odnosi na šemu
- Neophodno je biti pažljiv, pošto su podaci u eventima predstavljeni korišćenjem custom Debezium tipova podataka (naročito specijalizovani numerički tipovi i vremenski tipovi)

# PROJEKAT - DIAGRAM



# PROJEKT - DOCKER COMPOSE



Lulexs	added diagram	f3daea0 · 16 hours ago	2 Commits
redis_sink	first commit	20 hours ago	
README.md	added diagram	16 hours ago	
diagram.png	added diagram	16 hours ago	
docker-compose.yml	first commit	20 hours ago	
oracle-sink.json	first commit	20 hours ago	
oracle.sql	first commit	20 hours ago	
postgres-connector.json	first commit	20 hours ago	
postgres.sql	first commit	20 hours ago	

- 1 Kafka broker, koristi KRaft protokol
- Kafka UI, na portu 8070
- PostgreSQL
- DebeziumConnect
- Oracle
- Redis

```
lule@ASUS-LAPTOP:~/debezium_swe$ docker compose up -d
[+] Running 7/7
  ✓ Network debezium_swe_default  Created          0.0s
  ✓ Container kafka              Started         1.1s
  ✓ Container redis              Started         1.1s
  ✓ Container postgres            Started         1.2s
  ✓ Container oracle              Started         1.1s
  ✓ Container kafka-ui           Started         1.3s
  ✓ Container debezium-connect  Started         1.3s
```

# PROJEKT - POSTGRESQL SETUP



```
create table app_user (
    id int generated always as identity primary key,
    first_name varchar(50) not null,
    last_name varchar(50) not null,
    birthday date not null
);

create table user_attributes (
    user_id int not null references app_user(id),
    attribute_name varchar(50) not null,
    attribute_value varchar(50),
    primary key (user_id, attribute_name)
);

alter table app_user replica identity full;
alter table user_attributes replica identity full;
```

```
create role debezium with login password 'debezium' replication;
grant connect on database postgres to debezium;
grant usage on schema public to debezium;
grant select on table public.app_user, public.user_attributes to debezium;
create publication dbz_publication for table public.app_user, public.user_attributes;
alter publication dbz_publication owner to debezium;
select pg_create_logical_replication_slot(
    'debezium_slot',
    'pgoutput'
);
```

# PROJEKAT - DEBEZIUM SOURCE CONNECTOR



```
{  
  "name": "postgres-cdc",  
  "config": {  
    "connector.class": "io.debezium.connector.postgresql.PostgresConnector",  
    "database.hostname": "postgres",  
    "database.port": "5432",  
    "database.user": "debezium",  
    "database.password": "debezium",  
    "database dbname": "postgres",  
  
    "plugin.name": "pgoutput",  
    "slot.name": "debezium_slot",  
    "publication.name": "dbz_publication",  
  
    "snapshot.mode": "initial",  
  
    "key.converter": "org.apache.kafka.connect.json.JsonConverter",  
    "value.converter": "org.apache.kafka.connect.json.JsonConverter",  
  
    "topic.prefix": "dbz"  
  }  
}
```

```
lule@ASUS-LAPTOP:~/debezium_swe$ curl -X POST \  
  -H "Content-Type: application/json" \  
  --data @postgres-connector.json \  
  http://localhost:8083/connectors | jq  
% Total % Received % Xferd Average Speed Time Time Time Current  
Dload Upload Total Spent Left Speed  
100 1140 100 558 100 582 16005 16693 --:--:--:--:--:-- 33529  
{  
  "name": "postgres-cdc",  
  "config": {  
    "connector.class": "io.debezium.connector.postgresql.PostgresConnector",  
    "database.hostname": "postgres",  
    "database.port": "5432",  
    "database.user": "debezium",  
    "database.password": "debezium",  
    "database dbname": "postgres",  
    "plugin.name": "pgoutput",  
    "slot.name": "debezium_slot",  
    "publication.name": "dbz_publication",  
    "snapshot.mode": "initial",  
    "key.converter": "org.apache.kafka.connect.json.JsonConverter",  
    "value.converter": "org.apache.kafka.connect.json.JsonConverter",  
    "topic.prefix": "dbz",  
    "name": "postgres-cdc"  
  },  
  "tasks": [],  
  "type": "source"  
}
```

# PROJEKAT - DEBEZIUM SINK CONNECTOR



```
{  
  "name": "oracle-jdbc-sink",  
  "config": {  
    "connector.class": "io.debezium.connector.jdbc.JdbcSinkConnector",  
    "hibernate.dialect": "org.hibernate.dialect.OracleDialect",  
    "tasks.max": "1",  
  
    "connection.url": "jdbc:oracle:thin:@//oracle:1521/FREEPDB1",  
    "connection.username": "system",  
    "connection.password": "oracle",  
  
    "insert.mode": "upsert",  
    "delete.enabled": "true",  
  
    "primary.key.mode": "record_key",  
  
    "schema.evolution": "none",  
    "use.time.zone": "UTC",  
  
    "topics": "dbz.public.app_user,dbz.public.user_attributes",  
  
    "collection.name.format": "${source.table}",  
    "table.name.format": "${source.table}"  
  }  
}
```

```
lule@ASUS-LAPTOP:~/debezium_swe$ curl -X POST \  
-H "Content-Type: application/json" \  
--data @oracle-sink.json \  
http://localhost:8083/connectors | jq  
% Total    % Received % Xferd  Average Speed   Time     Time     Time  Current  
                                         Dload  Upload  Total  Spent  Left  Speed  
100  1259  100   616  100   643  19940  20814 --:--:-- --:--:-- --:--:-- 41966  
{  
  "name": "oracle-jdbc-sink",  
  "config": {  
    "connector.class": "io.debezium.connector.jdbc.JdbcSinkConnector",  
    "hibernate.dialect": "org.hibernate.dialect.OracleDialect",  
    "tasks.max": "1",  
    "connection.url": "jdbc:oracle:thin:@//oracle:1521/FREEPDB1",  
    "connection.username": "system",  
    "connection.password": "oracle",  
    "insert.mode": "upsert",  
    "delete.enabled": "true",  
    "primary.key.mode": "record_key",  
    "schema.evolution": "none",  
    "use.time.zone": "UTC",  
    "topics": "dbz.public.app_user,dbz.public.user_attributes",  
    "collection.name.format": "${source.table}",  
    "table.name.format": "${source.table}",  
    "name": "oracle-jdbc-sink"  
  },  
  "tasks": [],  
  "type": "sink"  
}
```

# PROJEKT - KAFKA EVENTS



```
insert into app_user(first_name, last_name, birthday)
|   |   |
values ('Luka', 'Velickovic', '03.11.2002.');
```

Offset	Partition	Timestamp	Key Preview	Value Preview
0	0	1/11/2026, 12:05:47.242	{"schema": {"type": "struct", "fields": [{"type": "int32", ...}}], "payload": {"before": null, "after": {"id": 1, "first_name": "Luka", "last_name": "Velickovic", "birthday": 11757}, "source": {"version": "2.7.3.Final", "connector": "postgresql", "name": "dbz", "ts_ms": 1768129546578, "snapshot": "false", "db": "postgres", "sequence": "[null, \"25109704\"]", "ts_us": 1768129546578904, "ts_ns": 1768129546578904000, "schema": "public", "table": "app_user", "txId": 768, "lsn": 25109704, "xmin": null}, "transaction": null}}	<p>Timestamp: 1/11/2026, 12:05:47.242 Timestamp type: CREATE_TIME</p> <p>Key Serde: String Size: 155 Bytes</p> <p>Value Serde: String Size: 3 KB</p>

# PROJEKT - KAFKA EVENTS



```
select * from app_user;
```

Script Output x | Query Result x

SQL | All Rows Fetched: 1 in 0.086 seconds

ID	FIRST_NAME	LAST_NAME	BIRTHDAY
1	Luka	Velickovic	11-MAR-02

HASH public-app\_user:1

104 B Length: 4 TTL: No limit

< 1 min ↗ Show TTL +

Field	Value	TTL
id	1	No Limit
first_name	Luka	No Limit
last_name	Velickovic	No Limit
birthday	11757	No Limit